


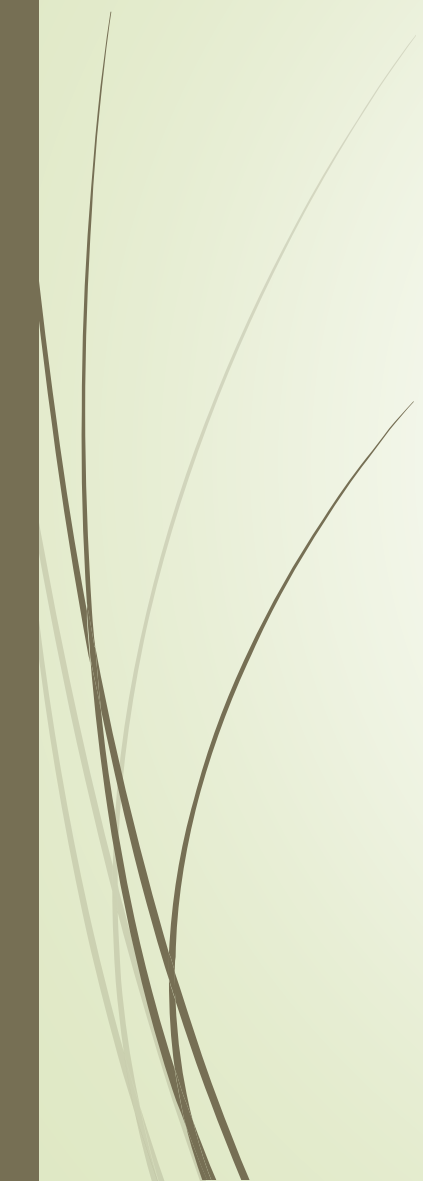
Binary Deep Learning



Presented by
Roey Nagar and Kostya Berestizshevsky



Lecture Outline

- Motivation and existing studies
 - BinaryConnect
 - XNOR-Net
 - Bonus: Hardware for machine learning
- 

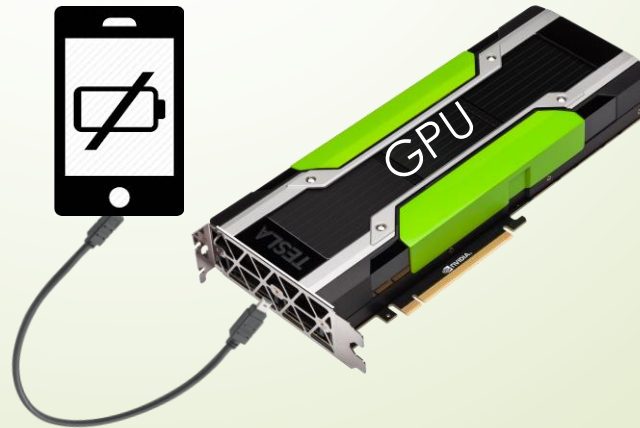
Motivation

- Make neural computation suitable for low-power **dedicated hardware**
- Reduce computational/memory effort (training / inference/ both)



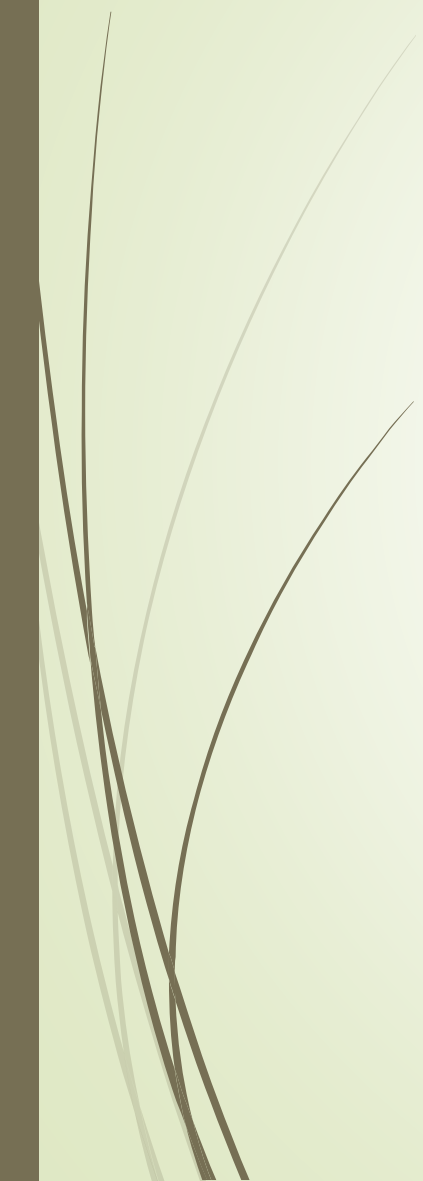
Alex Net:

- 61M parameters
- 1.5B high precision operations per image






Existing Methods To Reduce Over-parametrization

- Shallow-networks
 - Compress a pre-trained network
 - Compact Layers
 - Quantize parameters (to several levels)
 - Network binarization (to 2 values)
- 

BinaryConnect: Training Deep Neural Networks with binary weights during propagations

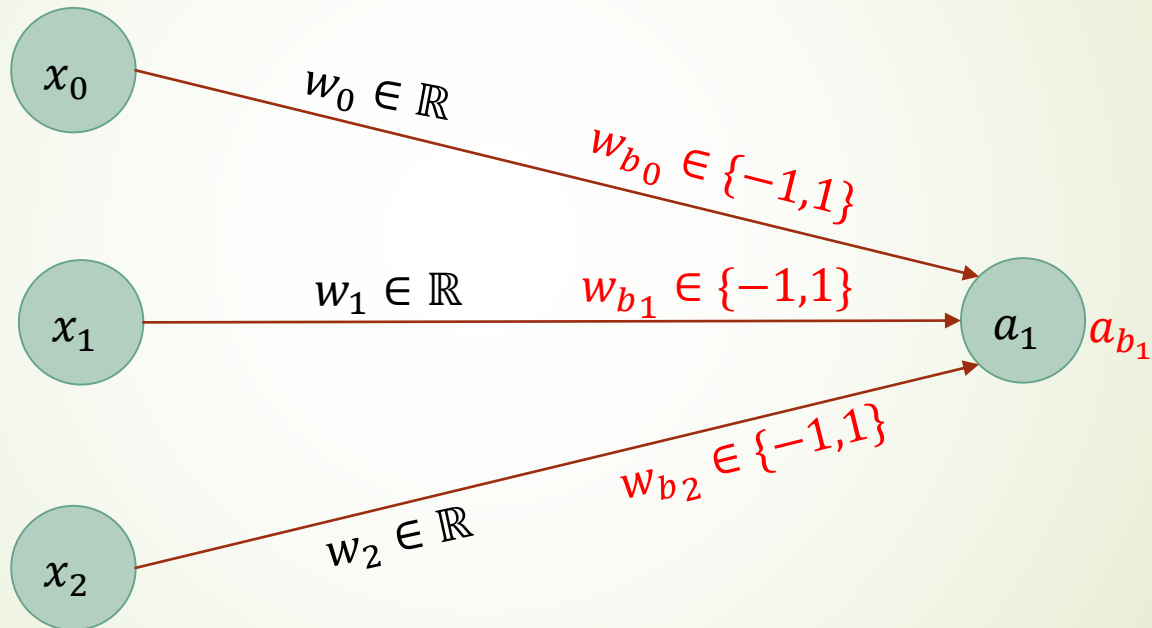


First presented at NIPS 2015 by
Matthieu Courbariaux
Yoshua Bengio
Jean-Pierre David

Today presented by
Roey Nagar

Idea

► Binarization of weights



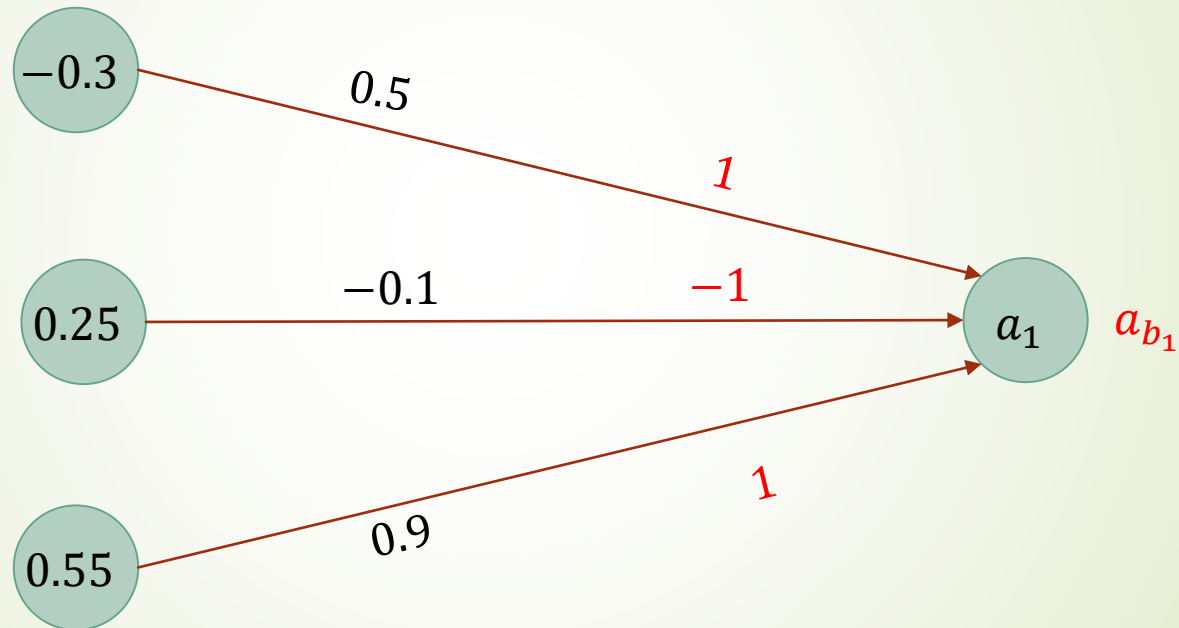
$$a_1 = \text{ReLU}(w_0 \cdot x_0 + w_1 \cdot x_1 + w_2 \cdot x_2)$$

$$a_{b_1} = \text{ReLU}(\underbrace{w_{b_0}}_{\text{Sign}} x_0 + \underbrace{w_{b_1}}_{\text{Sign}} x_1 + \underbrace{w_{b_2}}_{\text{Sign}} x_2)$$

No multiplications required

Idea

► Binarization of weights



$$a_1 = \text{ReLU}(0.5 \cdot (-0.3) - 0.1 \cdot 0.25 + 0.9 \cdot 0.55) = 0.32$$

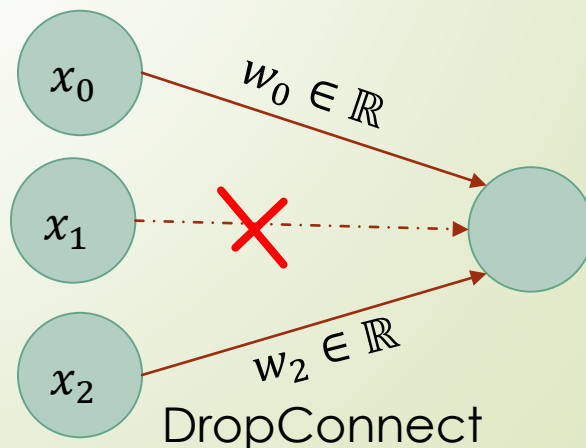
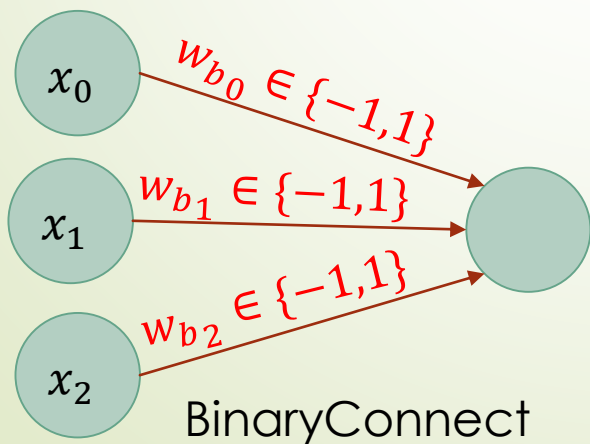
$$a_{b_1} = \text{ReLU}(-0.3 - 0.25 + 0.55) = 0$$

Ingredient #1 – Update w , not w_b

- ▶ $w_b \in \{-1, 1\}$ are obtained from w
- ▶ Forward/backward pass uses $w_b \in \{-1, 1\}$
- ▶ Only the real valued $w \in \mathbb{R}$ are updated
- ▶ **Important in order to catch the small contributions of the sub-gradients of SGD**

Ingredient #2 - Regularization

- Similarly to DropConnect, BinaryConnect alters the weights, but via binarizing them to $\{-1, 1\}$ instead of zeroing part of the weights each time.
- This can be referred to as a “quantization noise” or a **regularization**





BinaryConnect

In detail

Deterministic vs stochastic binarization

► Deterministic:

$$w_b = \begin{cases} +1 & \text{if } w \geq 0, \\ -1 & \text{otherwise} \end{cases}$$

► Stochastic

$$w_b = \begin{cases} +1 & \text{w.p. } p = \sigma(w), \\ -1 & \text{w.p. } 1 - p \end{cases}$$

$$\sigma(w) = \max\left(0, \min\left(1, \frac{x+1}{2}\right)\right)$$

BinaryConnect Training

Algorithm 1 SGD training with BinaryConnect. C is the cost function for minibatch and the functions $\text{binarize}(w)$ and $\text{clip}(w)$ specify how to binarize and clip weights. L is the number of layers.

Require: a minibatch of (inputs, targets), previous parameters w_{t-1} (weights) and b_{t-1} (biases), and learning rate η .

Ensure: updated parameters w_t and b_t .

1. Forward propagation:

$w_b \leftarrow \text{binarize}(w_{t-1})$

For $k = 1$ to L , compute a_k knowing a_{k-1} , w_b and b_{t-1}

2. Backward propagation:

Initialize output layer's activations gradient $\frac{\partial C}{\partial a_L}$

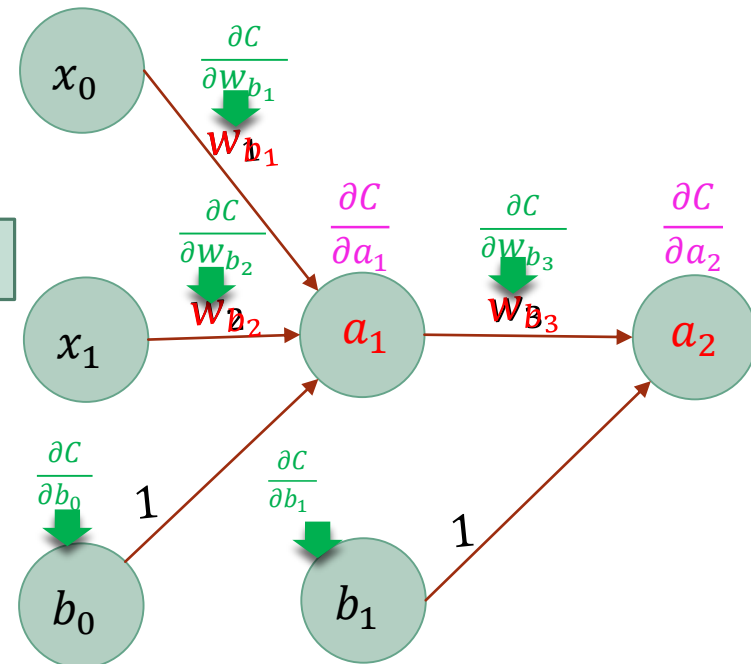
For $k = L$ to 2, compute $\frac{\partial C}{\partial a_{k-1}}$ knowing $\frac{\partial C}{\partial a_k}$ and w_b

3. Parameter update:

Compute $\frac{\partial C}{\partial w_b}$ and $\frac{\partial C}{\partial b_{t-1}}$ knowing $\frac{\partial C}{\partial a_k}$ and a_{k-1}

$w_t \leftarrow \text{clip}(w_{t-1} - \eta \frac{\partial C}{\partial w_b})$

$b_t \leftarrow b_{t-1} - \eta \frac{\partial C}{\partial b_{t-1}}$



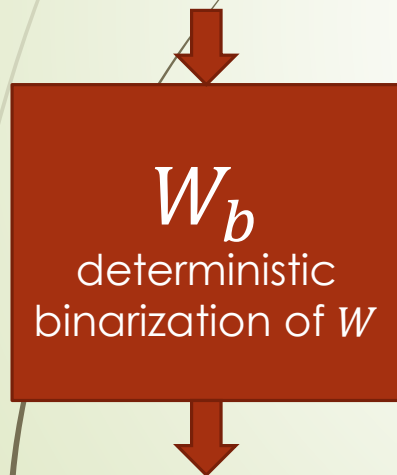
Training - Weight Clipping

- **Observation:** Binarization operation is not influenced by w with magnitude beyond 1.
- **Problem:** w might grow very large without influencing w_b
- **Solution:** The real valued w are clipped to $[-1, +1]$ upon updates
- This is yet another weight **regularization**

Test-Time Inference

- So far we trained DNN with on-the-fly binarization
- Result: we have the **trained real valued W**
- How do we test this DNN on **new examples**?

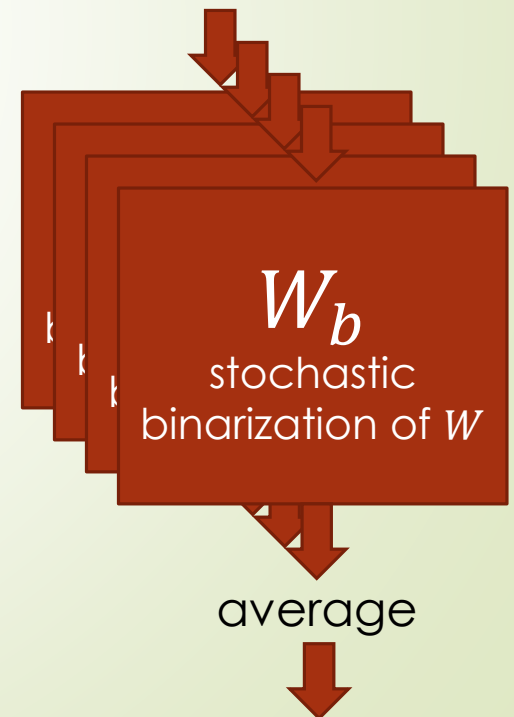
Option 1



Option 2



Option 3



Experiments

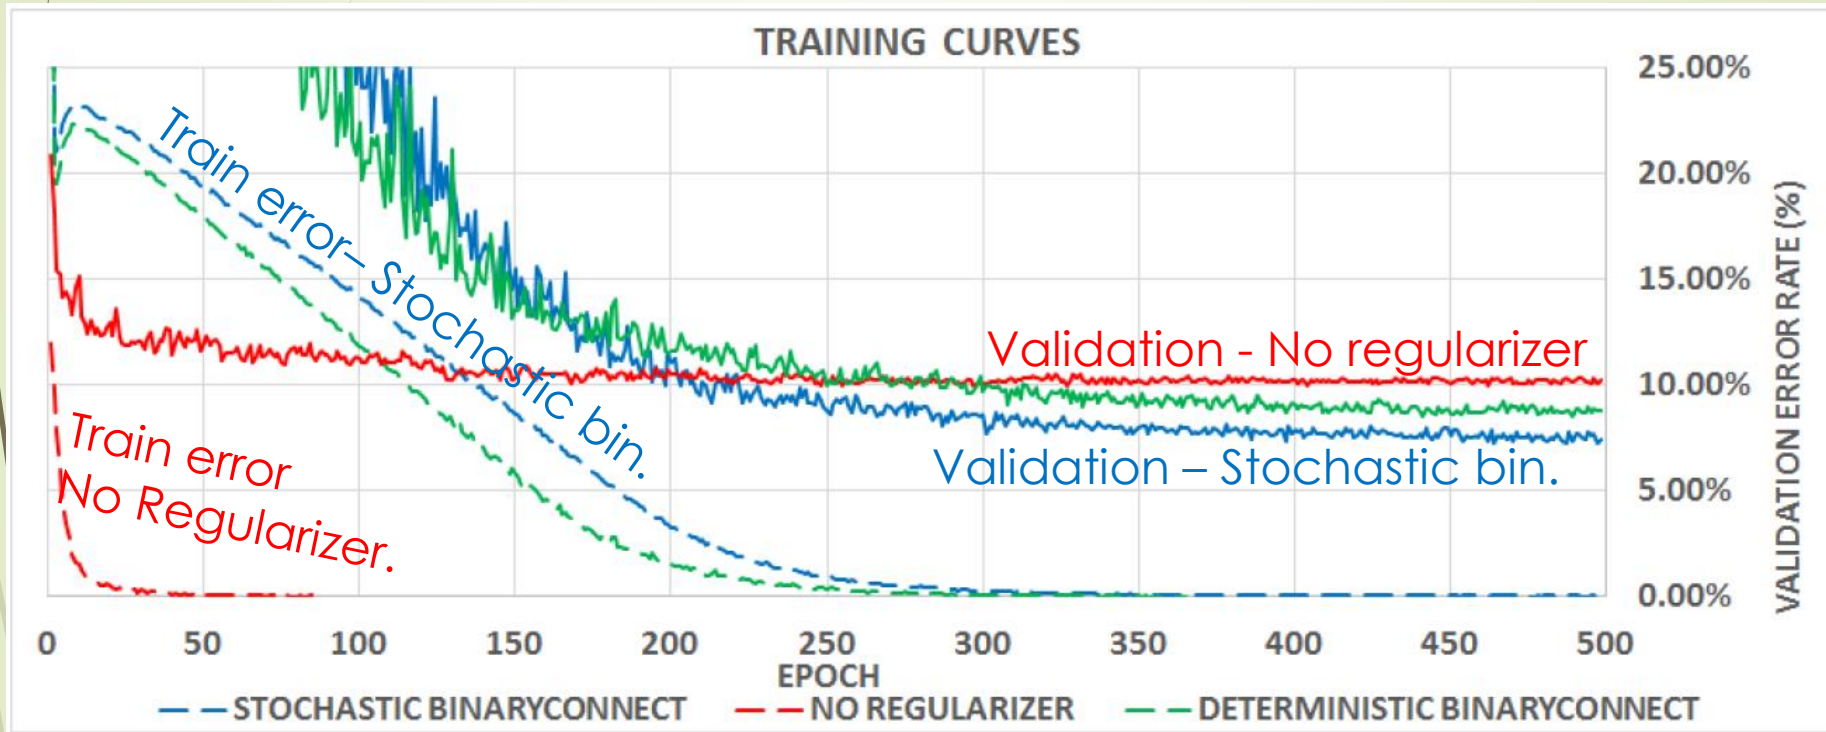
- **MNIST** – 3 FC layers + SVM
- **CIFAR-10, SVHN** – 6 CONV layers + 2 FC +SVM

Method	MNIST	CIFAR-10	SVHN
No regularizer	$1.30 \pm 0.04\%$	10.64%	2.44%
BinaryConnect (det.)	$1.29 \pm 0.08\%$	9.90%	2.30%
BinaryConnect (stoch.)	$1.18 \pm 0.04\%$	8.27%	2.15%
50% Dropout	$1.01 \pm 0.04\%$		
Maxout Networks [29]	0.94%	11.68%	2.47%
Deep L2-SVM [30]	0.87%		
Network in Network [31]		10.41%	2.35%
DropConnect [21]			1.94%
Deeply-Supervised Nets [32]		9.78%	1.92%

W_b

$W \in \mathbb{R}$

Regularization contribution

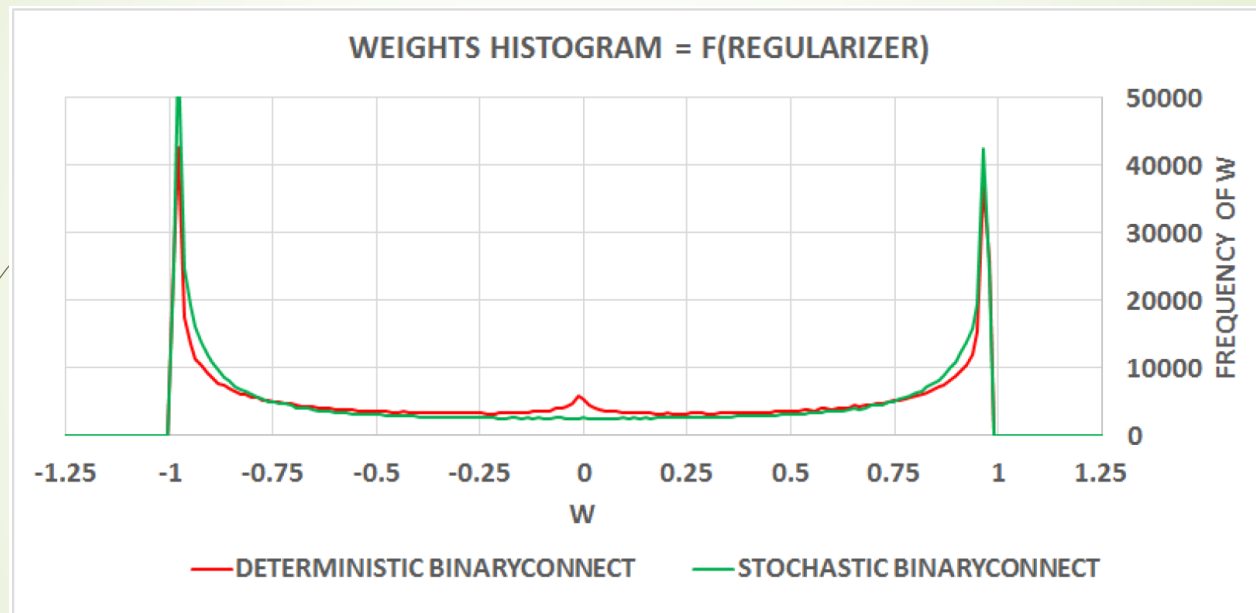


(CNN of CIFAR-10)

- Train Error and Time increases
- Validation Cost decreases
- Similar to the effect of Dropout

Training accomplished...

- How are the final weights distributed?



(FC network of MNIST)

- The weights are trying to become deterministic to reduce error

Similar Approaches

- **Ternary Connect** (SiPS 2014 Hwang et al.)
 - Train only using $w \in \mathbb{R}$
 - Ternarize the weights to $w_T \in \{-H, 0, H\}$ where H is adjusted to minimize error
 - Re-train using w_T during propagation and w for updates
- **BinaryNet (BNN)** (2016 Courbariaux et al.)
 - Binarize weights and activations

BinaryConnect Summary

- Weights are binarized during forward and backward propagations
- Acts as a regularizer, preserves accuracy
- Training - $\approx 2/3$ of the multiplications removed
- Test time - getting rid of the multiplications altogether
- Test time – memory bandwidth reduction $\times 16$ (16-bit floating point \rightarrow 1-bit)
- Theano + Lasagne code available:

<https://github.com/MatthieuCourbariaux/BinaryConnect>

XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks



First presented at ECCV 2016 by
Mohammad Rastegariy
Vicente Ordonezy
Joseph Redmon
Ali Farhadi

Today presented by
Kostya Berestizshevsky

From Google Ngram-Viewer





Highlights

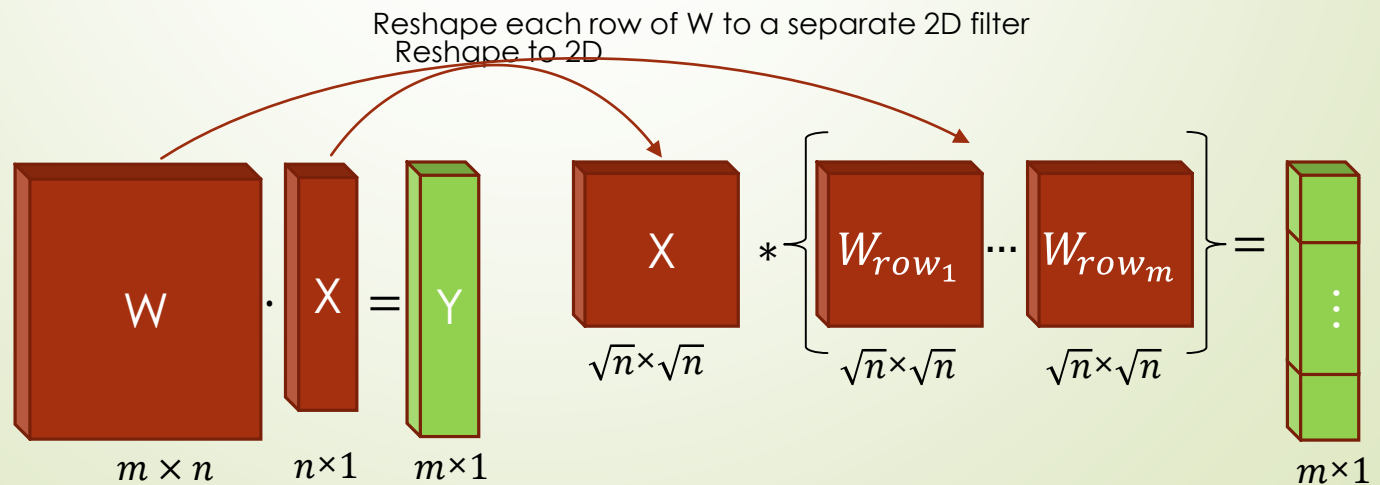
- Try binarization on **large datasets**
- Two new variations tested on ImageNet:
 1. Binary-Weight-Nets <bin weights only>
 2. XNOR-Nets <bin. weights & inputs>
- Lower run-times and memory
(relative to full-precision)
- Higher accuracy
(relative to former binary methods)

Interesting Observation

- FC layer with n inputs and m outputs (weight matrix of a size $m \times n$)



- CONV layer with m filters of a size n



Binary CNN 1st variation: Binary Weight Networks

- W - real valued weights (filter)
- I - real valued input tensor
- $*$ is a convolution operation

$$I * W \approx (I \oplus B)\alpha$$

- B – binary weights (the sign of W)
- \oplus is a convolution using only add/sub operations
- α – real valued scale factor (the average of $|W|$)

Why XNOR + bitcount approximates MAC ?

I	W	$I \odot W$	$XNOR(I, W)$
-1	-1	1	1
-1	1	-1	-1
1	-1	-1	-1
1	1	1	1

$I \cdot W$	$bitCount(XNOR(I, W))$
0	0

Binary CNN 2nd variation: XNOR Networks

- W - real valued weights (filter)
- I - real valued input tensor
- $*$ is a convolution operation

$$I * W \approx (H \circledast B) \odot \alpha K$$

- H – input binary tensor (the sign of I)
- B – binary weights (the sign of W)
- \circledast is a convolution using only XNOR and bit-counting
- α – real valued scale factors (the average of $|W|$)
- K – real valued scale factors of input conv-windows
- \odot is elementwise multiplication

XNOR-Net Convolution Example

I – input tensor

0.1	0.8	0.7	0.5
0.3	-0.9	-0.9	-0.8
0.5	-0.2	-0.4	-0.5
0.9	-0.4	1	0.8

W filter

-0.4	-0.5
1	0.8

$$\alpha = 0.225$$

$$K =$$

0.525	0.825	0.725
0.475	0.6	0.65
0.5	0.5	0.675

$I * W$

-0.86	-2.29	-2.07
0.67	0.29	-0.04
0.48	0.68	2.05

H – binarized input tensor

1	1	1	1
1	-1	-1	-1
1	-1	-1	-1
1	-1	1	1

B binarized filter

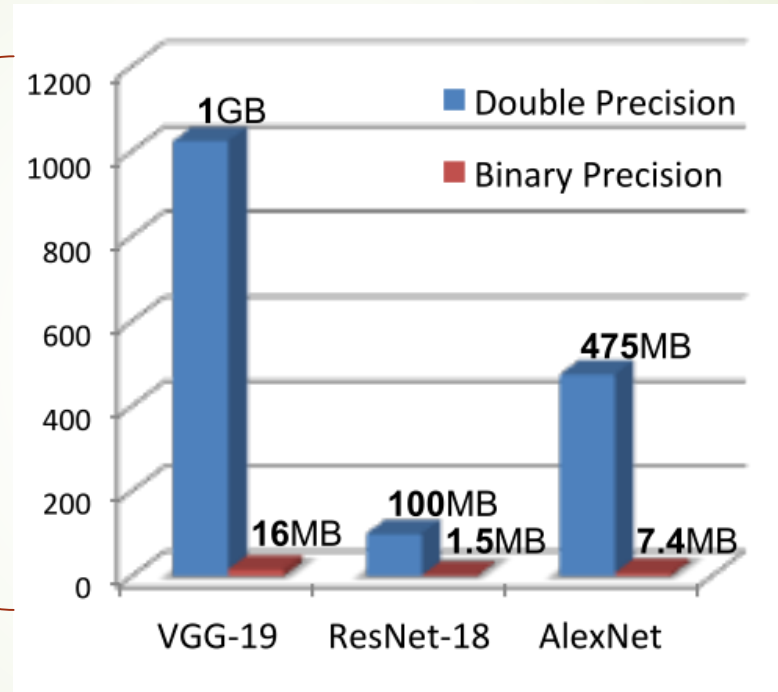
-1	-1
1	1

$H \odot B \odot K\alpha$

-0.71	-2.23	-1.96
0	0	0
0	0.68	1.82

Binary weights reduce memory

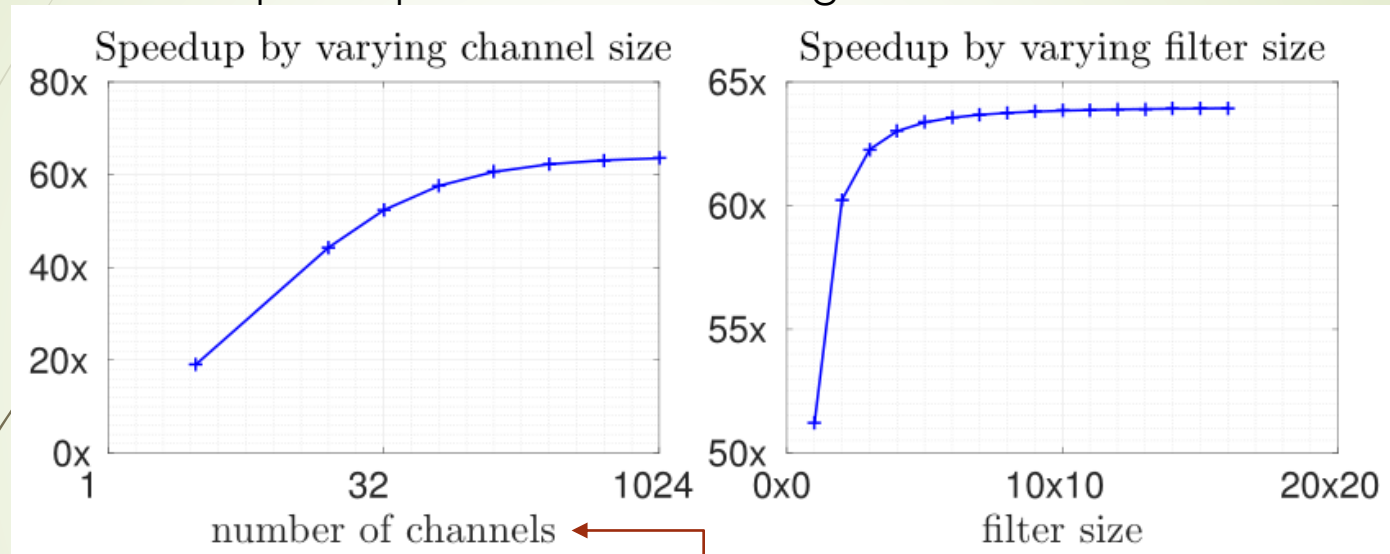
Required
memory
for weights
storage [MB]



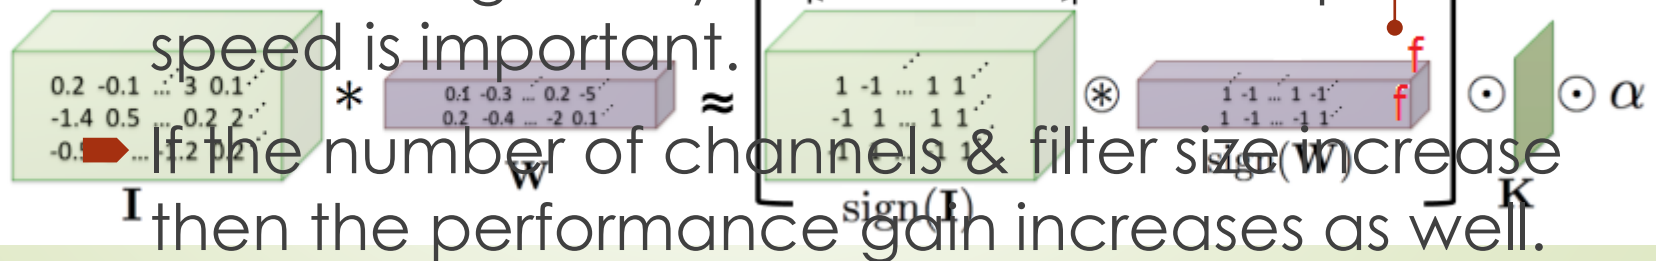
- Binary-weight-networks are so small that can be easily fitted into portable devices

Speedup using XNOR-bitcount

The speedup is relative to the regular convolution



Convolution with XNOR-Bitcount



Accuracy comparison

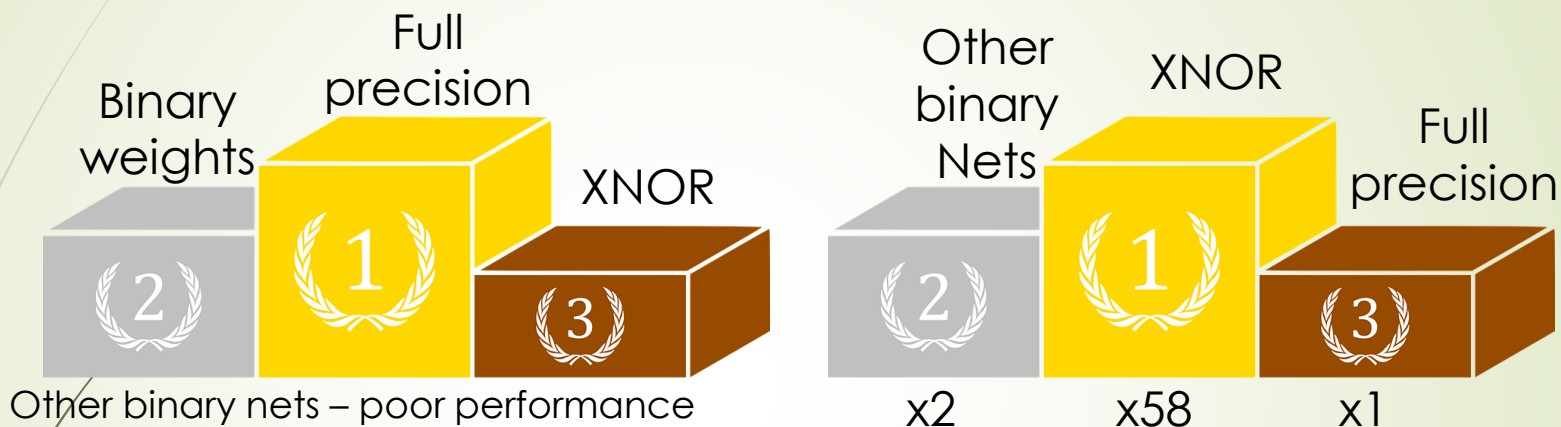
AlexNet Architecture

Classification Accuracy(%)									
Binary-Weight				Binary-Input-Binary-Weight				Full-Precision	
BWN		BC[11]		XNOR-Net		BNN[11]		AlexNet[1]	
Top-1	Top-5	Top-1	Top-5	Top-1	Top-5	Top-1	Top-5	Top-1	Top-5
56.8	79.4	35.4	61.0	44.2	69.2	27.9	50.42	56.6	80.2

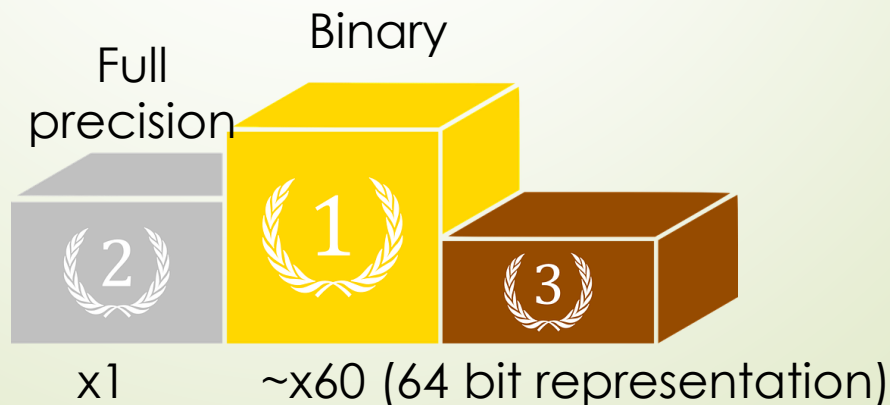
Other Architectures

	ResNet-18		GoogLenet	
Network Variations	top-1	top-5	top-1	top-5
Binary-Weight-Network	60.8	83.0	65.5	86.1
XNOR-Network	51.2	73.2	N/A	N/A
Full-Precision-Network	69.3	89.2	71.3	90.0

Tradeoff summary Speed

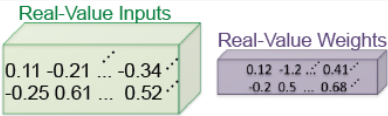
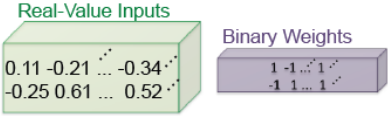
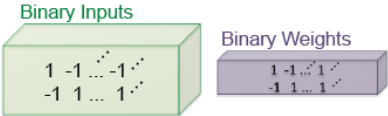


Memory



Conclusions

- Binary-Weight-Nets provide a good compromise of Memory & Speed vs. Accuracy
- XNOR-Nets provide supreme speedup at a cost of accuracy (The authors opened a startup!)

	Network Variations	Operations used in Convolution	Memory Saving (Inference)	Computation Saving (Inference)	Accuracy on ImageNet (AlexNet)
Standard Convolution		$+$, $-$, \times	1x	1x	%56.7
Binary Weight		$+$, $-$	$\sim 32x$	$\sim 2x$	%56.8
BinaryWeight Binary Input (XNOR-Net)		XNOR, bitcount	$\sim 32x$	$\sim 58x$	%44.2

- Torch 7 code available:
<https://github.com/allenai/XNOR-Net>



Dedicated Hardware for Machine Learning

(Bonus)



ML Hardware Research

- Dedicate hardware for training or inference?
- Optimize Throughput? [Gop/sec]
- Optimize Area? [Gop/sec/mm²]
- Optimize Power? [Gop/sec/Watt]
- ASIC / FPGA?

Intermediate Conclusions

- Memory is the bottleneck.
- Weights $\in \mathbb{R}$: 32-bits for training, 16-bits for infer.
- The design must be general enough – since the ML algorithms are changing rapidly,

Academic works

- **CNP'09, Neuflow'11, nn-X'14 – FPGA** [Farabet, Gokhale et al.]

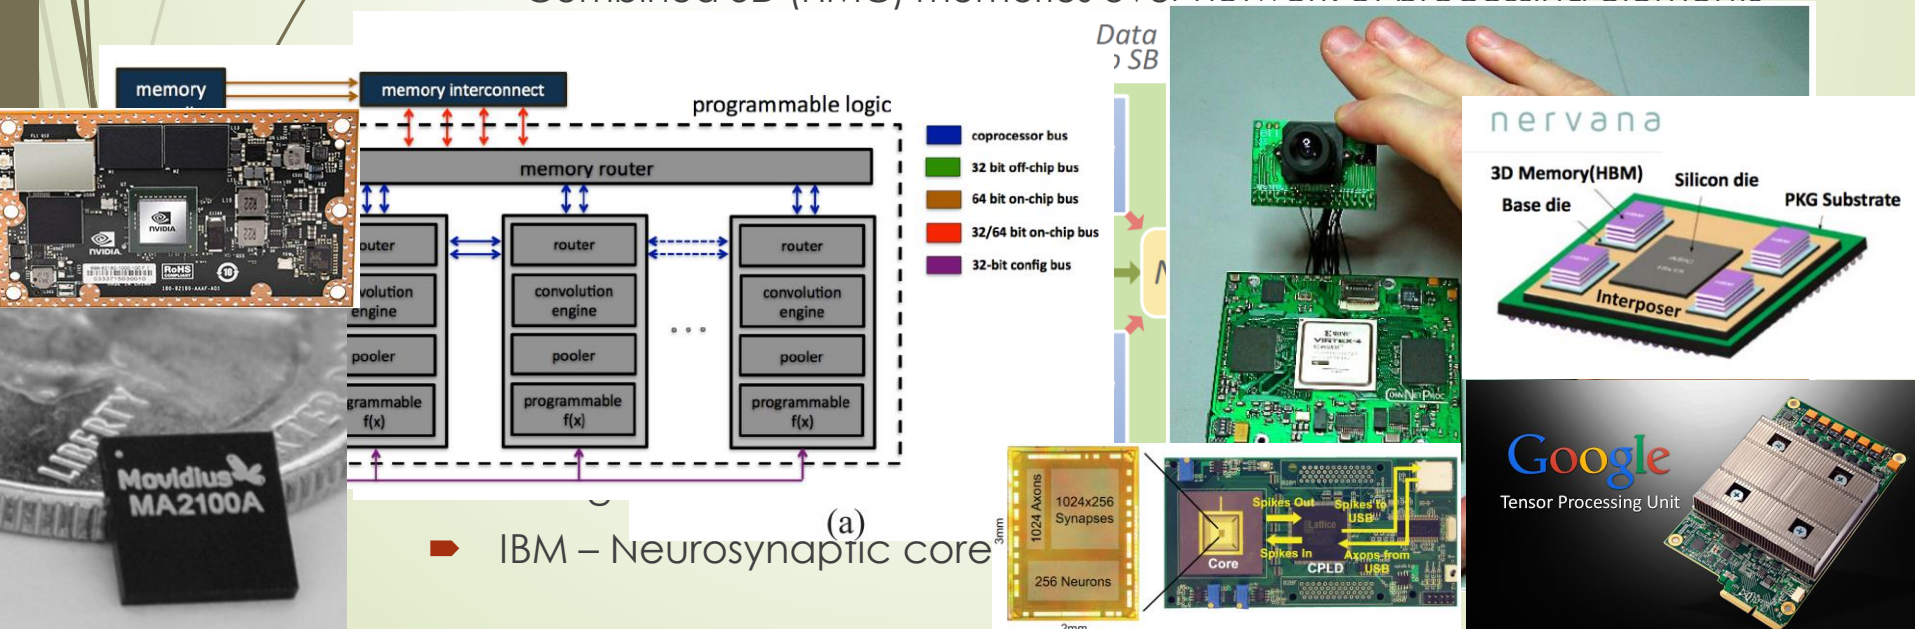
Low power of 4W yet high throughput of 200GOp/sec

- **DaDianNao – ASIC** [Chen, Temam et al.]

Focused on a smart memory placement around processing elements
x150 power reduction and x450 speedup relative to K20 GPU

- **Neurocube'16 – ASIC** [Kim et al.]

Combined 3D (HMC) memories over network of processing elements



Thank You!
Any Questions?



References – Binary DL

- [Matthieu Courbariaux](#), [Yoshua Bengio](#), [Jean-Pierre David](#): BinaryConnect: Training Deep Neural Networks with binary weights during propagations. [NIPS 2015](#): 3123-3131
- [Mohammad Rastegari](#), [Vicente Ordonez](#), [Joseph Redmon](#), [Ali Farhadi](#): XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks. [ECCV \(4\) 2016](#): 525-542
- [Matthieu Courbariaux](#), [Yoshua Bengio](#): BinaryNet: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1. [CoRRabs/1602.02830](#) (2016)
- K. Hwang and W. Sung, "Fixed-point feedforward deep neural network design using weights +1, 0, and -1," *2014 IEEE Workshop on Signal Processing Systems (SiPS)*, Belfast, 2014, pp. 1-6.

References - Hardware

- C. Farabet, C. Poulet, J. Y. Han and Y. LeCun, "**CNP**: An FPGA-based processor for Convolutional Networks," *2009 International Conference on Field Programmable Logic and Applications*, Prague, 2009, pp. 32-37.
- C. Farabet, B. Martini, B. Corda, P. Akselrod, E. Culurciello and Y. LeCun, "**NeuFlow**: A runtime reconfigurable dataflow processor for vision," *CVPR 2011 WORKSHOPS*, Colorado Springs, CO, 2011, pp. 109-116.
- Gokhale, Vinayak, et al. "A 240 g-ops/s mobile coprocessor for deep neural networks." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops*. 2014.
- Yunji Chen, Tao Luo, Shaoli Liu, Shijin Zhang, Liqiang He, Jia Wang, Ling Li, Tianshi Chen, Zhiwei Xu, Ninghui Sun, and Olivier Temam. 2014. **DaDianNao**: A Machine-Learning Supercomputer. In *Proceedings of the 47th Annual IEEE/ACM International Symposium on Microarchitecture*
- D. Kim, J. Kung, S. Chai, S. Yalamanchili and S. Mukhopadhyay, "**Neurocube**: A Programmable Digital Neuromorphic Architecture with High-Density 3D Memory," *2016 ACM/IEEE 43rd Annual International Symposium on Computer Architecture (ISCA)*, Seoul, 2016, pp. 380-392.